

# Neural\_Network\_EDA,\_Model\_and\_Prediction\_on\_Wine\_Data

April 30, 2023

## 1 Brief

Welcome to the EDA, Model and Prediction of wine data. One of the ways to determine wine quality is by its physicochemistry. The purpose of this study was to look at the physicochemical properties of the wine and analyze which model, of the ones we learned in the course, can give us the best fit, prediction and accuracy score of the quality of the wine.

## 2 Introduction

Ever wondered if a wine is good based on its physicochemical measurements? We will look into building a set of models to analyze exactly that!

In this project, the data consisted of 6497 observations across 11 physicochemical properties and the corresponding quality. The datasets were cleaned and explored. Then, we build a deep neural network and predict the quality of wine based on the model built.

After that, the accuracy test showed how well the model did at predicting the wine quality based on the physicochemical properties of the wine.

## 3 The Data

### 3.1 Sources

Data Set Source: > “Wine Quality Data Set.” UCI Machine Learning Repository, 7 Oct. 2009, <https://archive.ics.uci.edu/ml/datasets/Wine+Quality>

Data Set Research: > P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. Modeling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553, 2009.

Further Research on Wine properties: > Mor, Nuriel S., et al. “Wine Quality and Type Prediction from Physicochemical Properties Using Neural Networks for Machine Learning: A Free Software for Winemakers and Customers.” AgriRxiv, vol. 2022, 30 Jan. 2022, <https://doi.org/10.31220/agriRxiv.2022.00125>.

Further Research on Winemaking > “Winemaking.” Wikipedia: The Free Encyclopedia. Wikimedia Foundation, Inc, 22 July 2004, <https://en.wikipedia.org/wiki/Winemaking>. Accessed 7 Dec. 2022.

## 3.2 Description of Data

The data, acquired from UCI, is related with wine physicochemical inputs and sensory outputs for a Portuguese “Vinho Verde” wine.

### 3.2.1 Physicochemical Properties

In this dataset, the wine quality is determined by 11 physicochemical qualities: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, and alcohol. I did some outside research to understand what these qualities are and what their values mean.

**Fixed Acidity:** The fixed acids found in wines are tartaric, malic, citric, and succinic. Values are how much acid is in the wine based on the type of acid.

**Volatile acidity:** This is a measure of the low molecular weight of fatty acids.

**Critic Acid:** This is an organic acid, added to give wine its sour taste amongst other things.

**Residual Sugar:** The sugar that is leftover after the fermentation process.

**Chlorides:** This measures how much salt is in the wine.

**Free Sulfur Dioxide:** These are another type of salty additive generally used for oxidation. A fraction of this will react with sugars and fraction will not. This is the fraction that did not react.

**Total Sulfur Dioxide:** The total of reacted and free sulfur dioxide.

**Density:** The measure of mass per unit volume.

**pH:** The scale used to measure acidity and basicity of wine.

**Sulfites:** The chemical compounds that contain sulfite ions for preservation.

**Alcohol:** The Alcohol by Volume of wine.

### 3.2.2 Other Variables

Aside from the physicochemical properties, there are a couple other variables in this dataset.

**Quality:** This is a score or grade given to a wine in part based on the above physicochemical properties.

**Color:** Red wine or White wine.

## 3.3 Data Import

To start building models, the data must be first imported.

### 3.3.1 Import Python Libraries

Begin by importing all the libraries that will be used.

```
[26]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import math
import seaborn as sns
sns.set()

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn import preprocessing, model_selection
import sklearn.metrics as metrics
from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, \
    ↪classification_report

from tensorflow import keras
from tensorflow.keras import layers, optimizers, regularizers
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import Dense, Dropout, Input, Add, \
    ↪BatchNormalization, \
                                Softmax, Activation
from tensorflow.keras.optimizers import Adam

from keras.utils import plot_model
#from kt_utils import *
import keras.backend as K

# Ignore Future Warnings
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

### 3.3.2 Import Wine Quality Data

Next, There are 2 CSV files to import that identify the quality of wines: one for red wines and one for white wines. These will be concatenated and this forms our complete DataFrame.

Since 'color' is a string label, it can be encoded of using the label as a number so it can be used for analysis.

Finally, the column names are updated with underscores to enable easier coding.

```
[27]: # Import red wines
df_red = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/
↳wine-quality/winequality-red.csv', sep=';')
df_red['color'] = "red"

# # Import white wines
df_white = pd.read_csv('https://archive.ics.uci.edu/ml/
↳machine-learning-databases/wine-quality/winequality-white.csv', sep=';')
df_white['color'] = "white"

# # Combine red and white
df = pd.concat([df_red, df_white], ignore_index=True)
df.reindex()

# # Encode color label
df['color'] = df['color'].astype('category')
df['color_enc'] = df['color'].cat.codes

# Use underscores instead of spaces in column names
df.columns = [c.replace(' ', '_') for c in df.columns]

df
```

```
[27]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
...	...	...	...	...	...	
6492	6.2	0.21	0.29	1.6	0.039	
6493	6.6	0.32	0.36	8.0	0.047	
6494	6.5	0.24	0.19	1.2	0.041	
6495	5.5	0.29	0.30	1.1	0.022	
6496	6.0	0.21	0.38	0.8	0.020	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	
6492	24.0	92.0	0.99114	3.27	0.50	
6493	57.0	168.0	0.99490	3.15	0.46	
6494	30.0	111.0	0.99254	2.99	0.46	
6495	20.0	110.0	0.98869	3.34	0.38	

6496		22.0		98.0	0.98941	3.26	0.32
------	--	------	--	------	---------	------	------

	alcohol	quality	color	color_enc
0	9.4	5	red	0
1	9.8	5	red	0
2	9.8	5	red	0
3	9.8	6	red	0
4	9.4	5	red	0
...	...	...	...	...
6492	11.2	6	white	1
6493	9.6	5	white	1
6494	9.4	6	white	1
6495	12.8	7	white	1
6496	11.8	6	white	1

[6497 rows x 14 columns]

The data has been imported into a DataFrame and is now ready for cleaning and analysis.

## 4 Cleaning the Data and Exploratory Data Analysis

Conduct some exploratory data analysis to understand the data a little better.

First, the shape of the dataset.

```
[28]: print(f"Shape of Dataset: {df.shape}")
```

Shape of Dataset: (6497, 14)

To understand the spread of the data in the columns, look at the column ranges (min and max).

```
[29]: # Print min and max for each column
for c in df.columns:
    # color is a Categorical column and doesn't have a min and max
    if c == "color":
        continue

    print(f"{c:<22}: (min, max) = ({df[c].min():.2f}, {df[c].max():.2f})")
```

```
fixed_acidity      : (min, max) = (3.80, 15.90)
volatile_acidity   : (min, max) = (0.08, 1.58)
citric_acid        : (min, max) = (0.00, 1.66)
residual_sugar     : (min, max) = (0.60, 65.80)
chlorides          : (min, max) = (0.01, 0.61)
free_sulfur_dioxide : (min, max) = (1.00, 289.00)
total_sulfur_dioxide : (min, max) = (6.00, 440.00)
density            : (min, max) = (0.99, 1.04)
pH                 : (min, max) = (2.72, 4.01)
```

And then lets take a look at how many unique values are in each column.

```
[30]: fixed_acidity      106
      volatile_acidity  187
      citric_acid       89
      residual_sugar    316
      chlorides         214
      free_sulfur_dioxide 135
      total_sulfur_dioxide 276
      density          998
      pH              108
      sulphates       111
      alcohol         111
      quality         7
      color           2
      color_enc       2
      dtype: int64
```

```
[31]: a = [np.nan, None, [], {}, 'NaN', 'Null', 'NULL', 'None', 'NA', '?', '-', '.', '', '\n',
        ↪ ', ' ']
row_str = "{:<22}{:<15}{:<15}"
print(row_str.format("Column Name", "Real Nulls", "Null-like"))
print("-----")
for c in df.columns:
    string_null = np.array([x in a[2:] for x in df[c]])
    print(row_str.format(c, str(df[c].isnull().sum()), str(string_null.sum())))
```

6

pH	0	0
sulphates	0	0
alcohol	0	0
quality	0	0
color	0	0
color_enc	0	0

There are certain columns that do not affect the quality of wine. For example “color”. Whether a wine is white or red by itself has no bearing on its quality.

As a note sometimes the color of the grape does have a different average for a particular physico-chemical property than the other.

Therefore, before we start analyzing the data, the color column. the encoded color column can be used instead, allowing for some forms of analysis.

```
[32]: df = df.drop('color', axis=1)
df = df.drop('color_enc', axis=1)

df
```

```
[32]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
...	...	...	...	...	...	
6492	6.2	0.21	0.29	1.6	0.039	
6493	6.6	0.32	0.36	8.0	0.047	
6494	6.5	0.24	0.19	1.2	0.041	
6495	5.5	0.29	0.30	1.1	0.022	
6496	6.0	0.21	0.38	0.8	0.020	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	
6492	24.0	92.0	0.99114	3.27	0.50	
6493	57.0	168.0	0.99490	3.15	0.46	
6494	30.0	111.0	0.99254	2.99	0.46	
6495	20.0	110.0	0.98869	3.34	0.38	
6496	22.0	98.0	0.98941	3.26	0.32	

	alcohol	quality
0	9.4	5

1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...	...	...
6492	11.2	6
6493	9.6	5
6494	9.4	6
6495	12.8	7
6496	11.8	6

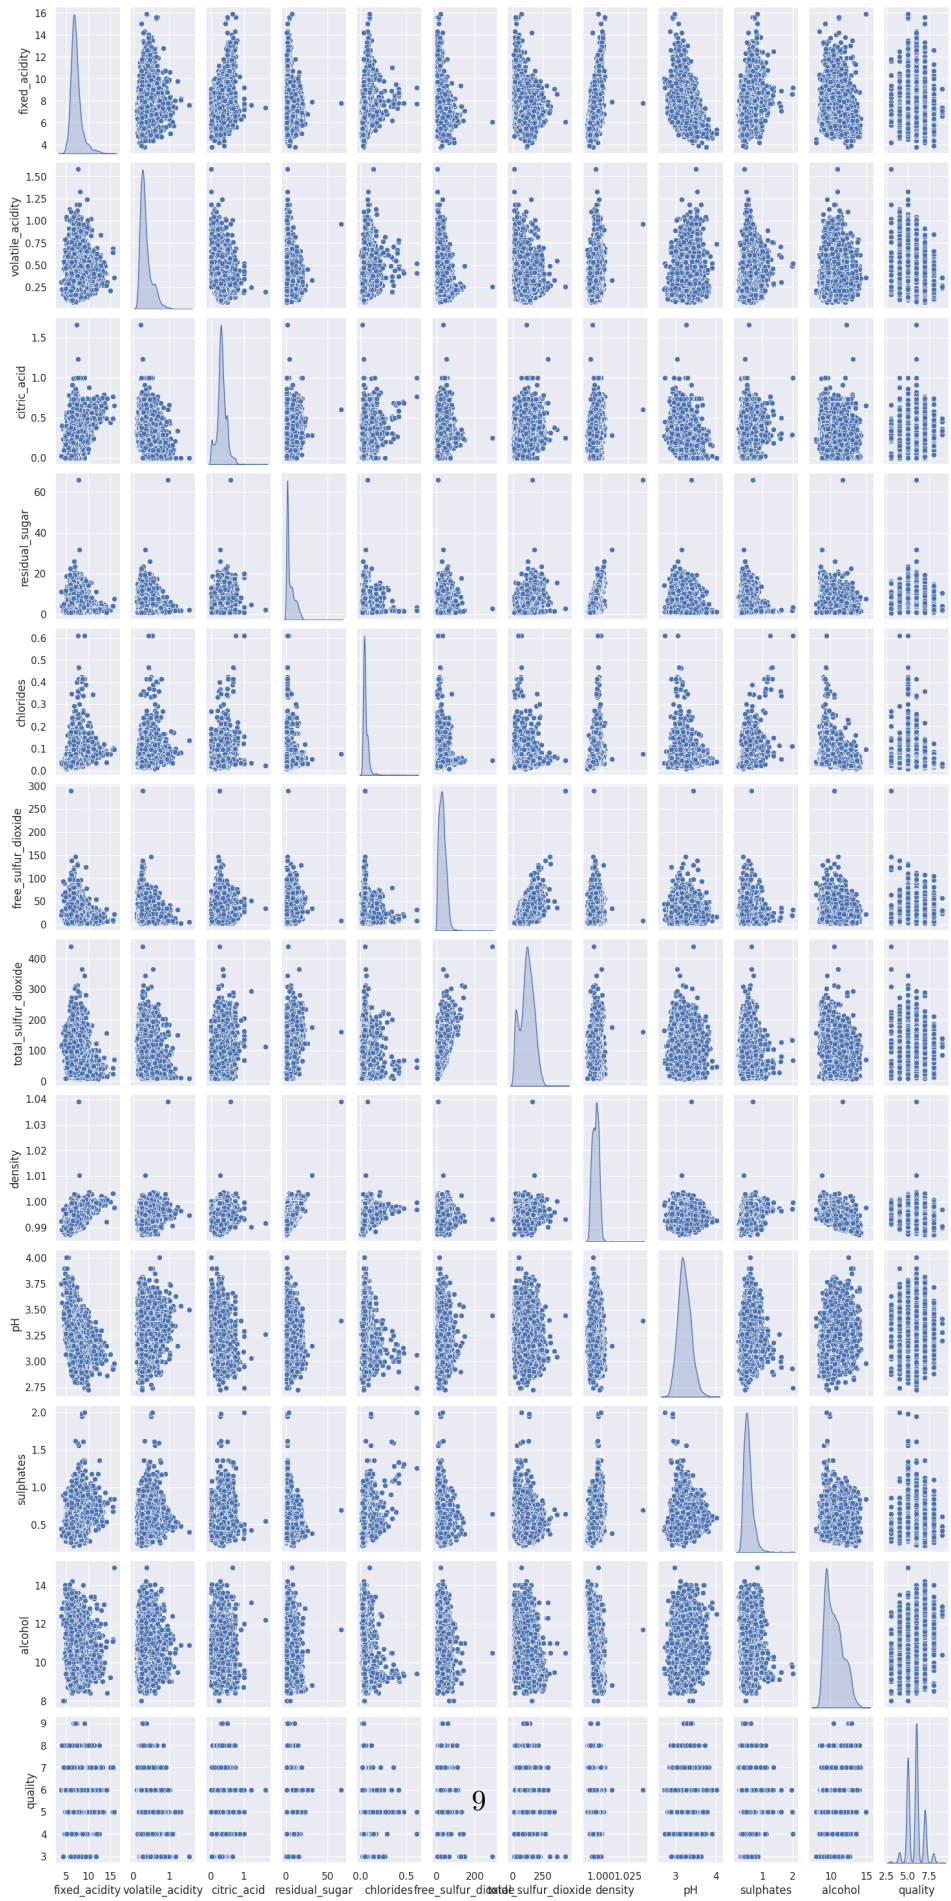
[6497 rows x 12 columns]

Inspecting the relationships between the features or components using a pair plot and a heatmap

```
[33]: # pair plot
sns.pairplot(df, aspect=0.5 , diag_kind='kde')
```

```
[33]: <seaborn.axisgrid.PairGrid at 0x7f57c79521a0>
```

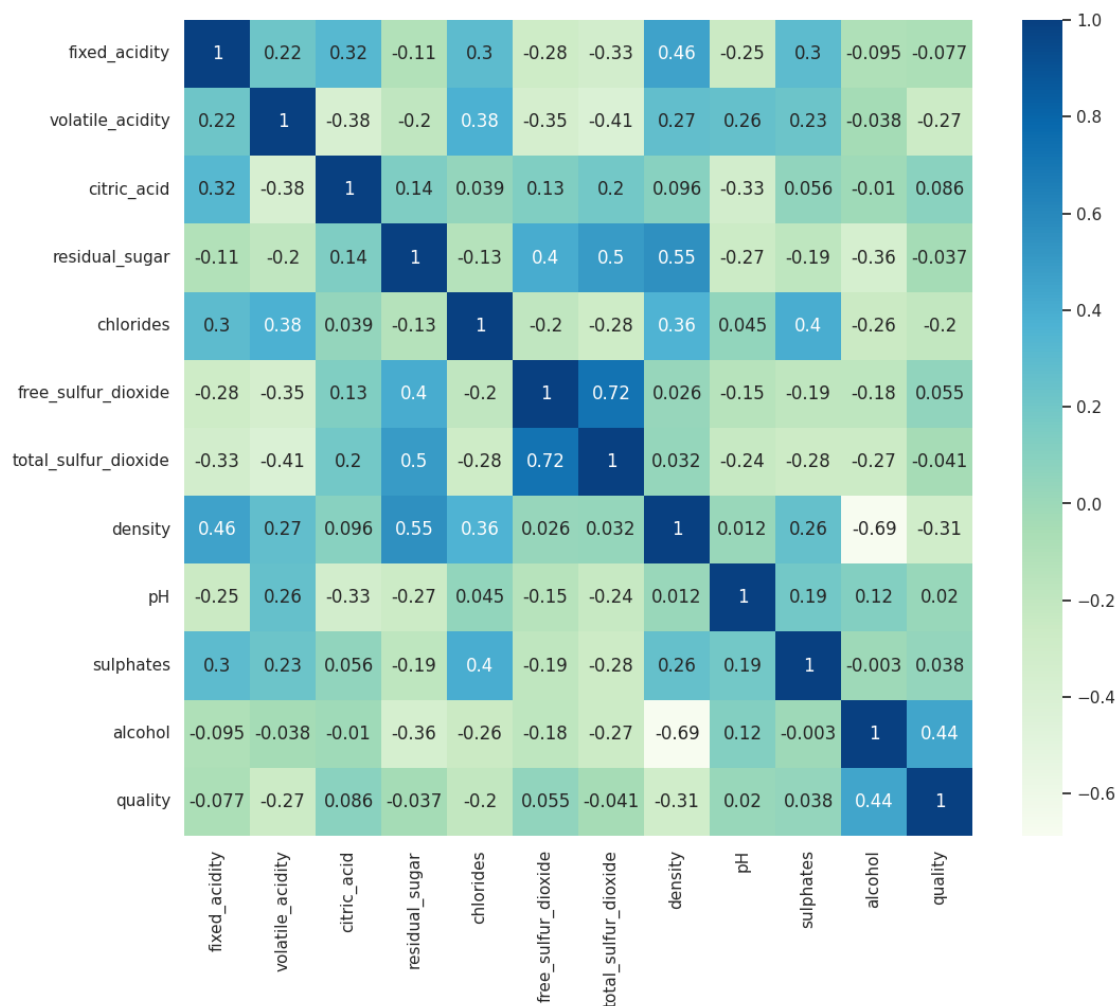




This is a histogram pairplot of quality vs each individual property. The chemical property is on the x-axis, while the quality is on the y-axis. This histogram pairplot shows the count of a particular rank at a particular value of the property. For example for fixed acidity, there are a of quality=6 wines in the range of 5 to 10 for fixed acidity vs, the range of 10 to 15. As the wine quality increases, more the lower range of the fixed acidity has more points than the upper range.

Next, plot a heatmap to visualize features are correlated with each other.

```
[34]: # Heatmap
plt.figure(figsize=(12,10))
cor = df.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.GnBu)
plt.show()
```

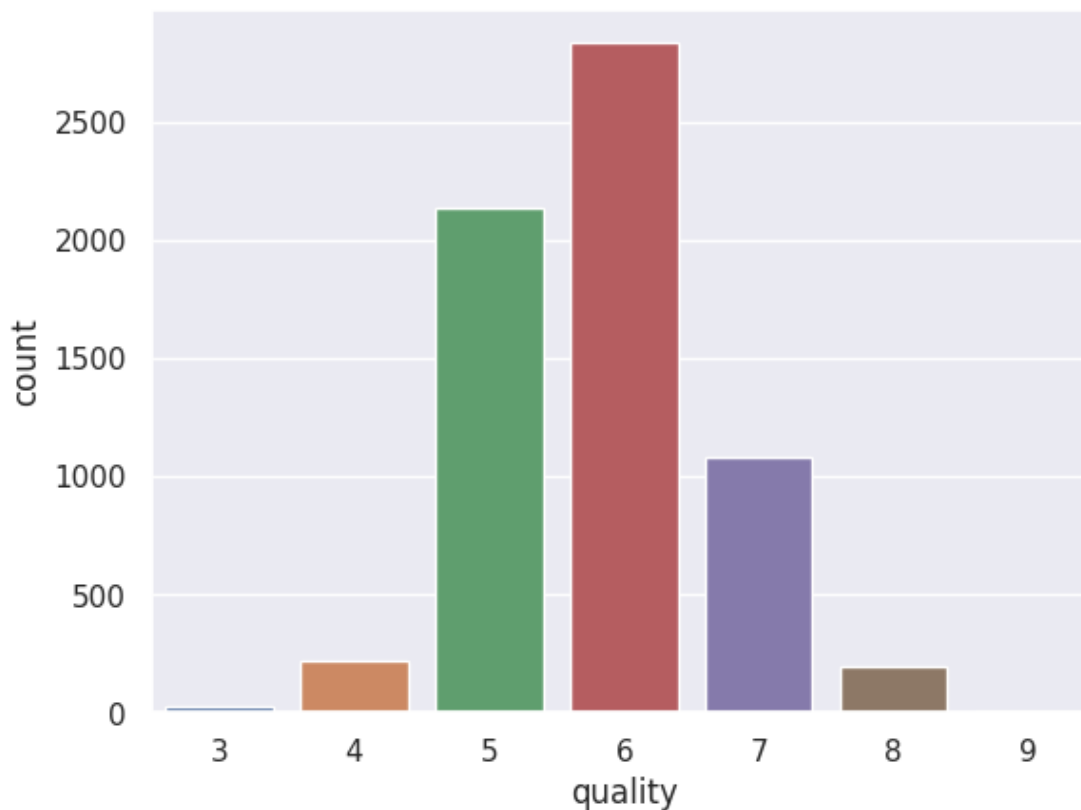


The Heat map shows the correlation between each of the properties. The highest correlation seems

to be between total sulfur dioxide and the color. This could mean that a particular color of wine has more total sulfur dioxide than the other color. For determining, which factor affects the quality the most, the highest correlation seems to be between quality and alcohol content.

```
[35]: sns.countplot(x=df["quality"])
```

```
[35]: <Axes: xlabel='quality', ylabel='count'>
```



Here is is distribution of the wines by quality. They are centered around 5, 6 and 7. This could be a downside for our model building.

```
[36]: df["quality"] =df["quality"].astype(int)
df = pd.get_dummies(df, columns=["quality"])

df.head(5)
```

```
[36]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality_3	quality_4	quality_5	quality_6	quality_7	quality_8	\
0	9.4	0	0	1	0	0	0	
1	9.8	0	0	1	0	0	0	
2	9.8	0	0	1	0	0	0	
3	9.8	0	0	0	1	0	0	
4	9.4	0	0	1	0	0	0	

	quality_9
0	0
1	0
2	0
3	0
4	0

## 5 Modeling the Data

### 5.1 Preparing the Data

First we prepare the data so that we do not have bias later. We will also split the data into training and test set. The training set will be 75% and the test set will be 25%.

```
[37]: X = df.iloc[:,0:11].values # first columns
      Y = df.iloc[:,12:].values # last columns

      X = preprocessing.normalize(X, axis = 0)

      X_train,X_test,Y_train,Y_test = model_selection.
      ↪train_test_split(X,Y,test_size=0.25)

      print(X_train.shape,Y_train.shape,X_test.shape,Y_test.shape)
```

```
(4872, 11) (4872, 6) (1625, 11) (1625, 6)
```

### 5.2 Method

1. **Evaluate Models** Build a model

2. **Make Predictions or Plot the Accuracy** Make predictions and find accuracy or metrics of the model.
3. **Analyze** Analyze and make conclusion about the fit and accuracy of the model.

### 5.2.1 Model Architecture

With a stronger sense of the training data and the goal, the next discussion is which model architecture to use and train for best results. Selecting the right model architecture is important because it will use the right amount of resources for the task, will not overfit the data, and will give accurate classification results.

### 5.2.2 Options for Architecture

The loss function must be selected, with possible choices including mean absolute error (mae), mean squared error (mse), and Huber loss for regression loss calculations. This problem is a categorical classification problem so there are the options of using Categorical Cross Entropy.

Additionally the activation function for the linear units must be selected. Choices here include Sigmoid, TanH, and ReLU. Since ReLU has shown great promise with training Neural Network models and since it also can increase the speed of the model learning, that will be one of the activation function selected here.

Additionally, it is possible to select strategies to optimize the hyperparameters of an Neural Network model and one of the most popular and effective strategies is Stochastic Gradient Descent. This will be the selected strategy for optimizing the Neural Network to reduce loss.

```
[38]: def model_performance_graphs(classifier, acc):

    fig, axes = plt.subplots(1, 2, figsize = (8, 3))

    axes[0].plot(classifier.epoch, classifier.history[acc], label = 'acc')
    axes[0].set_title('Accuracy vs Epochs', fontsize = 12)
    axes[0].set_xlabel('Epochs', fontsize = 10)
    axes[0].set_ylabel('Accuracy', fontsize = 10)
    axes[0].legend()

    axes[1].plot(classifier.epoch, classifier.history['loss'], label = 'loss')
    axes[1].set_title("Loss Curve",fontsize=12)
    axes[1].set_xlabel("Epochs",fontsize=10)
    axes[1].set_ylabel("Loss",fontsize=10)
    axes[1].legend()

    plt.show()
```

### 5.3 Basic Stochastic Gradient Descent

Stochastic gradient descent (SGD) is a method for finding a minimum of a function by iteratively taking steps in the direction of the negative gradient of the function at the current point. The steps are chosen randomly, hence the name “stochastic”. SGD is a popular algorithm for training neural networks.

SGD works by repeatedly taking steps in the direction of the negative gradient of the loss function. The loss function is a measure of how well the model fits the training data. The gradient of the loss function is a vector that points in the direction of greatest increase in the loss function. By taking steps in the opposite direction of the gradient, SGD can find a minimum of the loss function.

```
[39]: mod1 = keras.Sequential([
        layers.Dense(512, activation='relu', input_shape=[11]),
        layers.Dense(512, activation='relu'),
        layers.Dense(512, activation='relu'),
        layers.Dense(1),
    ])
```

```
[40]: mod1.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	6144
dense_5 (Dense)	(None, 512)	262656
dense_6 (Dense)	(None, 512)	262656
dense_7 (Dense)	(None, 1)	513

```
=====
Total params: 531,969
Trainable params: 531,969
Non-trainable params: 0
=====
```

```
[41]: mod1.compile(
        optimizer='adam',
        loss='mae',
        metrics = ["accuracy"]
    )
```

```
[42]: history1 = mod1.fit(
        X_train, Y_train,
        validation_data=(X_test, Y_test),
```

```
    batch_size=256,  
    epochs=100,  
)
```

```
Epoch 1/100  
20/20 [=====] - 2s 37ms/step - loss: 0.1703 - accuracy:  
0.8340 - val_loss: 0.1664 - val_accuracy: 0.8344  
Epoch 2/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1666 - accuracy:  
0.8340 - val_loss: 0.1658 - val_accuracy: 0.8344  
Epoch 3/100  
20/20 [=====] - 1s 43ms/step - loss: 0.1661 - accuracy:  
0.8340 - val_loss: 0.1656 - val_accuracy: 0.8344  
Epoch 4/100  
20/20 [=====] - 1s 45ms/step - loss: 0.1664 - accuracy:  
0.8340 - val_loss: 0.1664 - val_accuracy: 0.8344  
Epoch 5/100  
20/20 [=====] - 1s 47ms/step - loss: 0.1665 - accuracy:  
0.8340 - val_loss: 0.1662 - val_accuracy: 0.8344  
Epoch 6/100  
20/20 [=====] - 1s 36ms/step - loss: 0.1664 - accuracy:  
0.8340 - val_loss: 0.1659 - val_accuracy: 0.8344  
Epoch 7/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val_loss: 0.1657 - val_accuracy: 0.8344  
Epoch 8/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1663 - accuracy:  
0.8340 - val_loss: 0.1659 - val_accuracy: 0.8344  
Epoch 9/100  
20/20 [=====] - 1s 29ms/step - loss: 0.1663 - accuracy:  
0.8340 - val_loss: 0.1660 - val_accuracy: 0.8344  
Epoch 10/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1662 - accuracy:  
0.8340 - val_loss: 0.1660 - val_accuracy: 0.8344  
Epoch 11/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1662 - accuracy:  
0.8340 - val_loss: 0.1659 - val_accuracy: 0.8344  
Epoch 12/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1662 - accuracy:  
0.8340 - val_loss: 0.1659 - val_accuracy: 0.8344  
Epoch 13/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1662 - accuracy:  
0.8340 - val_loss: 0.1658 - val_accuracy: 0.8344  
Epoch 14/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val_loss: 0.1658 - val_accuracy: 0.8344  
Epoch 15/100
```

20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 16/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1662 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 17/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1663 - accuracy:  
0.8340 - val\_loss: 0.1661 - val\_accuracy: 0.8344  
Epoch 18/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1666 - accuracy:  
0.8340 - val\_loss: 0.1662 - val\_accuracy: 0.8344  
Epoch 19/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1665 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 20/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 21/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 22/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 23/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 24/100  
20/20 [=====] - 1s 33ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 25/100  
20/20 [=====] - 1s 43ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1659 - val\_accuracy: 0.8344  
Epoch 26/100  
20/20 [=====] - 1s 44ms/step - loss: 0.1662 - accuracy:  
0.8340 - val\_loss: 0.1661 - val\_accuracy: 0.8344  
Epoch 27/100  
20/20 [=====] - 1s 47ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 28/100  
20/20 [=====] - 1s 32ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 29/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 30/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1660 - val\_accuracy: 0.8344  
Epoch 31/100



20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 32/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1662 - accuracy:  
0.8340 - val\_loss: 0.1660 - val\_accuracy: 0.8344  
Epoch 33/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 34/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 35/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 36/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 37/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1662 - accuracy:  
0.8340 - val\_loss: 0.1656 - val\_accuracy: 0.8344  
Epoch 38/100  
20/20 [=====] - 1s 25ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1659 - val\_accuracy: 0.8344  
Epoch 39/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 40/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 41/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 42/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 43/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 44/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 45/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 46/100  
20/20 [=====] - 1s 31ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 47/100

20/20 [=====] - 1s 43ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 48/100  
20/20 [=====] - 1s 44ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 49/100  
20/20 [=====] - 1s 48ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 50/100  
20/20 [=====] - 1s 30ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 51/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 52/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 53/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 54/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 55/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 56/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 57/100  
20/20 [=====] - 1s 25ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 58/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 59/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 60/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1659 - val\_accuracy: 0.8344  
Epoch 61/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 62/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 63/100

20/20 [=====] - 1s 28ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 64/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 65/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 66/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 67/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 68/100  
20/20 [=====] - 1s 35ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 69/100  
20/20 [=====] - 1s 43ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 70/100  
20/20 [=====] - 1s 48ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 71/100  
20/20 [=====] - 1s 48ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 72/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 73/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 74/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 75/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 76/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 77/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 78/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 79/100

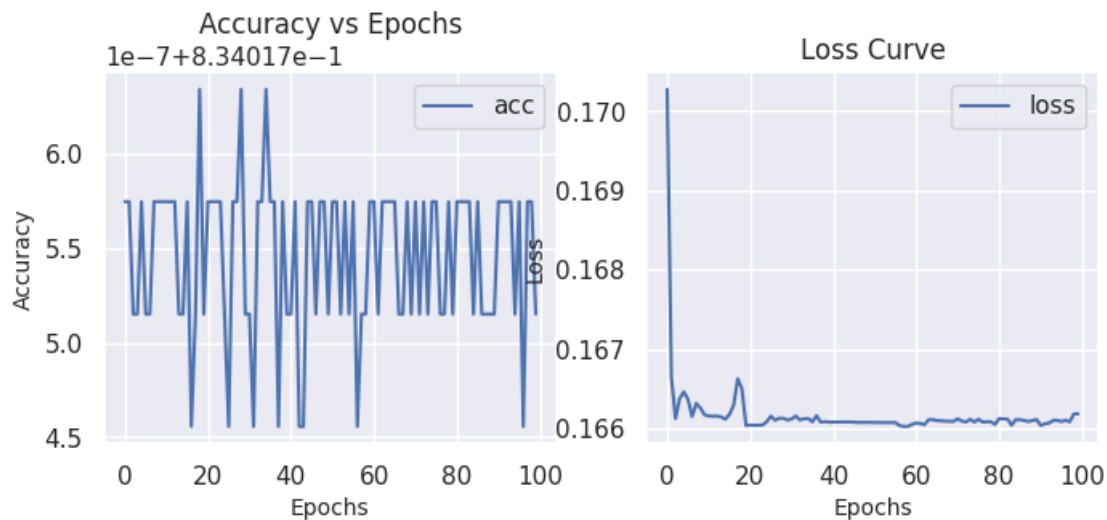
20/20 [=====] - 1s 27ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 80/100  
20/20 [=====] - 1s 29ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 81/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 82/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 83/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 84/100  
20/20 [=====] - 1s 27ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 85/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 86/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 87/100  
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 88/100  
20/20 [=====] - 1s 25ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1656 - val\_accuracy: 0.8344  
Epoch 89/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 90/100  
20/20 [=====] - 1s 39ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1656 - val\_accuracy: 0.8344  
Epoch 91/100  
20/20 [=====] - 1s 46ms/step - loss: 0.1660 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 92/100  
20/20 [=====] - 1s 47ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1659 - val\_accuracy: 0.8344  
Epoch 93/100  
20/20 [=====] - 1s 39ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1658 - val\_accuracy: 0.8344  
Epoch 94/100  
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:  
0.8340 - val\_loss: 0.1657 - val\_accuracy: 0.8344  
Epoch 95/100

```

20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:
0.8340 - val_loss: 0.1656 - val_accuracy: 0.8344
Epoch 96/100
20/20 [=====] - 1s 26ms/step - loss: 0.1661 - accuracy:
0.8340 - val_loss: 0.1657 - val_accuracy: 0.8344
Epoch 97/100
20/20 [=====] - 1s 28ms/step - loss: 0.1661 - accuracy:
0.8340 - val_loss: 0.1657 - val_accuracy: 0.8344
Epoch 98/100
20/20 [=====] - 0s 25ms/step - loss: 0.1661 - accuracy:
0.8340 - val_loss: 0.1657 - val_accuracy: 0.8344
Epoch 99/100
20/20 [=====] - 1s 30ms/step - loss: 0.1662 - accuracy:
0.8340 - val_loss: 0.1658 - val_accuracy: 0.8344
Epoch 100/100
20/20 [=====] - 1s 26ms/step - loss: 0.1662 - accuracy:
0.8340 - val_loss: 0.1657 - val_accuracy: 0.8344

```

```
[43]: model_performance_graphs(history1, "accuracy")
```



```

[44]: train1_acc = mod1.evaluate(x = X_train, y = Y_train)
print()
print ("Train Accuracy = " + str(train1_acc[1]))
print()
test1_acc = mod1.evaluate(x = X_test, y = Y_test)
print()
print ("Loss = " + str(test1_acc[0]))
print ("Test Accuracy = " + str(test1_acc[1]))

```

```

153/153 [=====] - 1s 3ms/step - loss: 0.1660 -

```

```
accuracy: 0.8340
```

```
Train Accuracy = 0.8340182900428772
```

```
51/51 [=====] - 0s 4ms/step - loss: 0.1657 - accuracy: 0.8344
```

```
Loss = 0.165687695145607
```

```
Test Accuracy = 0.8343589305877686
```

The accuracy of our model is pretty high. We can attempt to improve the accuracy by utilizing other techniques

## 5.4 Dropout and Batch Normalization

Dropout and batch normalization are two techniques used to improve the performance of neural networks. Dropout works by randomly dropping out (setting to zero) a certain percentage of neurons during each training epoch. This forces the network to learn more robust features that are not dependent on any single neuron. Batch normalization works by normalizing the inputs to each layer of the network, which helps to prevent the values from becoming too large or too small.

```
[45]: mod2 = Sequential()
      # layer 1
      mod2.add(Dense(30, input_dim=11, activation='relu',
      ↪name='fc0',kernel_regularizer=regularizers.l2(0.01)))
      mod2.add(BatchNormalization(momentum=0.99, epsilon=0.001))
      #layer 2
      mod2.add(Dense(50, name='fc1',bias_initializer='zeros'))
      mod2.add(Activation('tanh'))
      mod2.add(BatchNormalization(momentum=0.99, epsilon=0.001))
      mod2.add(Dropout(0.5))
      #layer 3
      mod2.add(Dense(100, name='fc2',bias_initializer='zeros'))
      mod2.add(Activation('relu'))
      mod2.add(BatchNormalization(momentum=0.99, epsilon=0.001))
      mod2.add(Dropout(0.5))
      #layer 4
      mod2.add(Dense(6, name='fc3',bias_initializer='zeros'))
      mod2.add(Activation('softmax'))
      mod2.add(BatchNormalization(momentum=0.99, epsilon=0.001))
```

```
[46]: mod2.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
fc0 (Dense)	(None, 30)	360

batch_normalization_4 (Batch Normalization)	(None, 30)	120
fc1 (Dense)	(None, 50)	1550
activation_3 (Activation)	(None, 50)	0
batch_normalization_5 (Batch Normalization)	(None, 50)	200
dropout_2 (Dropout)	(None, 50)	0
fc2 (Dense)	(None, 100)	5100
activation_4 (Activation)	(None, 100)	0
batch_normalization_6 (Batch Normalization)	(None, 100)	400
dropout_3 (Dropout)	(None, 100)	0
fc3 (Dense)	(None, 6)	606
activation_5 (Activation)	(None, 6)	0
batch_normalization_7 (Batch Normalization)	(None, 6)	24

```

=====
Total params: 8,360
Trainable params: 7,988
Non-trainable params: 372
-----

```

```

[47]: Adam = optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999,
    ↪epsilon=1e-08)
mod2.compile(optimizer = Adam, loss = "mae", metrics = ["accuracy"])

```

```

[48]: history2 = mod2.fit(x = X_train, y = Y_train, epochs = 100, verbose=1,
    ↪batch_size = 64, validation_data=(X_test, Y_test))

```

```

Epoch 1/100
77/77 [=====] - 3s 8ms/step - loss: 0.9232 - accuracy:
0.1882 - val_loss: 0.3982 - val_accuracy: 0.1631
Epoch 2/100
77/77 [=====] - 0s 5ms/step - loss: 0.7745 - accuracy:
0.2377 - val_loss: 0.3616 - val_accuracy: 0.1729

```

Epoch 3/100  
77/77 [=====] - 0s 5ms/step - loss: 0.6275 - accuracy:  
0.2917 - val\_loss: 0.3279 - val\_accuracy: 0.4215  
Epoch 4/100  
77/77 [=====] - 0s 5ms/step - loss: 0.4973 - accuracy:  
0.3502 - val\_loss: 0.3253 - val\_accuracy: 0.4215  
Epoch 5/100  
77/77 [=====] - 0s 5ms/step - loss: 0.4119 - accuracy:  
0.3890 - val\_loss: 0.2613 - val\_accuracy: 0.4215  
Epoch 6/100  
77/77 [=====] - 0s 5ms/step - loss: 0.3449 - accuracy:  
0.4438 - val\_loss: 0.2515 - val\_accuracy: 0.4289  
Epoch 7/100  
77/77 [=====] - 0s 5ms/step - loss: 0.2929 - accuracy:  
0.4516 - val\_loss: 0.2518 - val\_accuracy: 0.4843  
Epoch 8/100  
77/77 [=====] - 0s 5ms/step - loss: 0.2496 - accuracy:  
0.4645 - val\_loss: 0.2365 - val\_accuracy: 0.4720  
Epoch 9/100  
77/77 [=====] - 0s 5ms/step - loss: 0.2230 - accuracy:  
0.4750 - val\_loss: 0.2075 - val\_accuracy: 0.5058  
Epoch 10/100  
77/77 [=====] - 1s 9ms/step - loss: 0.2091 - accuracy:  
0.4873 - val\_loss: 0.2112 - val\_accuracy: 0.4751  
Epoch 11/100  
77/77 [=====] - 1s 7ms/step - loss: 0.2003 - accuracy:  
0.4971 - val\_loss: 0.2494 - val\_accuracy: 0.4185  
Epoch 12/100  
77/77 [=====] - 1s 7ms/step - loss: 0.1957 - accuracy:  
0.4984 - val\_loss: 0.2180 - val\_accuracy: 0.4412  
Epoch 13/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1879 - accuracy:  
0.5082 - val\_loss: 0.1904 - val\_accuracy: 0.4929  
Epoch 14/100  
77/77 [=====] - 1s 7ms/step - loss: 0.1846 - accuracy:  
0.5064 - val\_loss: 0.1846 - val\_accuracy: 0.4862  
Epoch 15/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1797 - accuracy:  
0.5101 - val\_loss: 0.1884 - val\_accuracy: 0.4609  
Epoch 16/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1775 - accuracy:  
0.5080 - val\_loss: 0.1873 - val\_accuracy: 0.4695  
Epoch 17/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1750 - accuracy:  
0.5113 - val\_loss: 0.1799 - val\_accuracy: 0.4837  
Epoch 18/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1716 - accuracy:  
0.5226 - val\_loss: 0.1684 - val\_accuracy: 0.5182



Epoch 19/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1706 - accuracy:  
0.5205 - val\_loss: 0.1763 - val\_accuracy: 0.4917  
Epoch 20/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1687 - accuracy:  
0.5261 - val\_loss: 0.1622 - val\_accuracy: 0.5329  
Epoch 21/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1697 - accuracy:  
0.5168 - val\_loss: 0.1657 - val\_accuracy: 0.5255  
Epoch 22/100  
77/77 [=====] - 1s 7ms/step - loss: 0.1655 - accuracy:  
0.5287 - val\_loss: 0.1687 - val\_accuracy: 0.5108  
Epoch 23/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1665 - accuracy:  
0.5220 - val\_loss: 0.1686 - val\_accuracy: 0.5175  
Epoch 24/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1672 - accuracy:  
0.5174 - val\_loss: 0.1839 - val\_accuracy: 0.4597  
Epoch 25/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1643 - accuracy:  
0.5255 - val\_loss: 0.1599 - val\_accuracy: 0.5274  
Epoch 26/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1631 - accuracy:  
0.5277 - val\_loss: 0.1902 - val\_accuracy: 0.2849  
Epoch 27/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1642 - accuracy:  
0.5244 - val\_loss: 0.1879 - val\_accuracy: 0.2154  
Epoch 28/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1632 - accuracy:  
0.5271 - val\_loss: 0.1669 - val\_accuracy: 0.5108  
Epoch 29/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1646 - accuracy:  
0.5189 - val\_loss: 0.1599 - val\_accuracy: 0.5286  
Epoch 30/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1592 - accuracy:  
0.5308 - val\_loss: 0.1855 - val\_accuracy: 0.4505  
Epoch 31/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1624 - accuracy:  
0.5086 - val\_loss: 0.1647 - val\_accuracy: 0.5040  
Epoch 32/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1600 - accuracy:  
0.5216 - val\_loss: 0.1697 - val\_accuracy: 0.4228  
Epoch 33/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1611 - accuracy:  
0.5183 - val\_loss: 0.1700 - val\_accuracy: 0.4234  
Epoch 34/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1602 - accuracy:  
0.5193 - val\_loss: 0.1698 - val\_accuracy: 0.4209

Epoch 35/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1617 - accuracy: 0.5129 - val\_loss: 0.1691 - val\_accuracy: 0.4332  
Epoch 36/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1625 - accuracy: 0.5109 - val\_loss: 0.1556 - val\_accuracy: 0.5440  
Epoch 37/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1609 - accuracy: 0.5170 - val\_loss: 0.1706 - val\_accuracy: 0.4295  
Epoch 38/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1599 - accuracy: 0.5226 - val\_loss: 0.1758 - val\_accuracy: 0.4215  
Epoch 39/100  
77/77 [=====] - 1s 9ms/step - loss: 0.1597 - accuracy: 0.5195 - val\_loss: 0.1692 - val\_accuracy: 0.3520  
Epoch 40/100  
77/77 [=====] - 1s 9ms/step - loss: 0.1606 - accuracy: 0.5148 - val\_loss: 0.1746 - val\_accuracy: 0.4178  
Epoch 41/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1607 - accuracy: 0.5150 - val\_loss: 0.1659 - val\_accuracy: 0.4732  
Epoch 42/100  
77/77 [=====] - 1s 7ms/step - loss: 0.1609 - accuracy: 0.5148 - val\_loss: 0.1563 - val\_accuracy: 0.5372  
Epoch 43/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1592 - accuracy: 0.5236 - val\_loss: 0.1684 - val\_accuracy: 0.5028  
Epoch 44/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1595 - accuracy: 0.5181 - val\_loss: 0.1556 - val\_accuracy: 0.5311  
Epoch 45/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1585 - accuracy: 0.5209 - val\_loss: 0.1695 - val\_accuracy: 0.4283  
Epoch 46/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1597 - accuracy: 0.5135 - val\_loss: 0.1703 - val\_accuracy: 0.4166  
Epoch 47/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1616 - accuracy: 0.5041 - val\_loss: 0.1690 - val\_accuracy: 0.4222  
Epoch 48/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1595 - accuracy: 0.5156 - val\_loss: 0.1686 - val\_accuracy: 0.4437  
Epoch 49/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1584 - accuracy: 0.5222 - val\_loss: 0.1654 - val\_accuracy: 0.4615  
Epoch 50/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1593 - accuracy: 0.5113 - val\_loss: 0.1614 - val\_accuracy: 0.5200

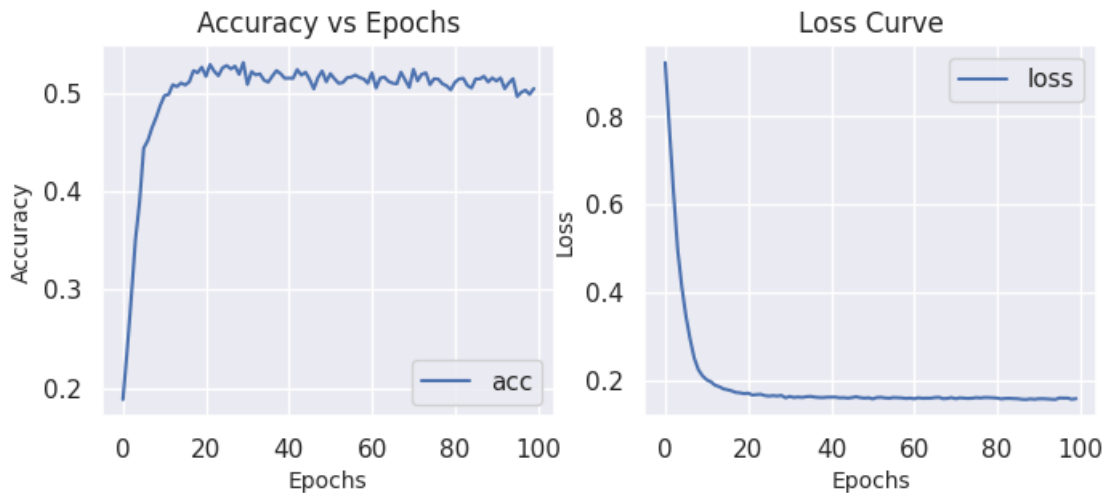
Epoch 51/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1567 - accuracy:  
0.5195 - val\_loss: 0.1802 - val\_accuracy: 0.4597  
Epoch 52/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1599 - accuracy:  
0.5142 - val\_loss: 0.1984 - val\_accuracy: 0.4037  
Epoch 53/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1603 - accuracy:  
0.5096 - val\_loss: 0.1634 - val\_accuracy: 0.5083  
Epoch 54/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1588 - accuracy:  
0.5105 - val\_loss: 0.1673 - val\_accuracy: 0.4997  
Epoch 55/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1582 - accuracy:  
0.5154 - val\_loss: 0.1553 - val\_accuracy: 0.5397  
Epoch 56/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1598 - accuracy:  
0.5160 - val\_loss: 0.1672 - val\_accuracy: 0.5034  
Epoch 57/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1593 - accuracy:  
0.5181 - val\_loss: 0.1557 - val\_accuracy: 0.5323  
Epoch 58/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1589 - accuracy:  
0.5156 - val\_loss: 0.1686 - val\_accuracy: 0.4295  
Epoch 59/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1576 - accuracy:  
0.5142 - val\_loss: 0.1605 - val\_accuracy: 0.5225  
Epoch 60/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1585 - accuracy:  
0.5099 - val\_loss: 0.1583 - val\_accuracy: 0.5182  
Epoch 61/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1570 - accuracy:  
0.5203 - val\_loss: 0.1555 - val\_accuracy: 0.5262  
Epoch 62/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1592 - accuracy:  
0.5049 - val\_loss: 0.1863 - val\_accuracy: 0.4418  
Epoch 63/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1580 - accuracy:  
0.5150 - val\_loss: 0.1596 - val\_accuracy: 0.5286  
Epoch 64/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1583 - accuracy:  
0.5160 - val\_loss: 0.1528 - val\_accuracy: 0.5391  
Epoch 65/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1583 - accuracy:  
0.5117 - val\_loss: 0.1652 - val\_accuracy: 0.4812  
Epoch 66/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1598 - accuracy:  
0.5099 - val\_loss: 0.1609 - val\_accuracy: 0.4991

Epoch 67/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1603 - accuracy: 0.5092 - val\_loss: 0.1690 - val\_accuracy: 0.3458  
Epoch 68/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1571 - accuracy: 0.5209 - val\_loss: 0.1569 - val\_accuracy: 0.5385  
Epoch 69/100  
77/77 [=====] - 1s 9ms/step - loss: 0.1578 - accuracy: 0.5131 - val\_loss: 0.1559 - val\_accuracy: 0.5372  
Epoch 70/100  
77/77 [=====] - 1s 7ms/step - loss: 0.1596 - accuracy: 0.5055 - val\_loss: 0.1888 - val\_accuracy: 0.4369  
Epoch 71/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1570 - accuracy: 0.5238 - val\_loss: 0.1523 - val\_accuracy: 0.5397  
Epoch 72/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1591 - accuracy: 0.5123 - val\_loss: 0.1531 - val\_accuracy: 0.5446  
Epoch 73/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1584 - accuracy: 0.5189 - val\_loss: 0.1554 - val\_accuracy: 0.5194  
Epoch 74/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1577 - accuracy: 0.5203 - val\_loss: 0.1684 - val\_accuracy: 0.4985  
Epoch 75/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1583 - accuracy: 0.5084 - val\_loss: 0.1600 - val\_accuracy: 0.5132  
Epoch 76/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1595 - accuracy: 0.5144 - val\_loss: 0.1699 - val\_accuracy: 0.4794  
Epoch 77/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1588 - accuracy: 0.5140 - val\_loss: 0.1576 - val\_accuracy: 0.5372  
Epoch 78/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1599 - accuracy: 0.5092 - val\_loss: 0.1674 - val\_accuracy: 0.4978  
Epoch 79/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1597 - accuracy: 0.5072 - val\_loss: 0.1623 - val\_accuracy: 0.5243  
Epoch 80/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1593 - accuracy: 0.5031 - val\_loss: 0.1565 - val\_accuracy: 0.5323  
Epoch 81/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1585 - accuracy: 0.5107 - val\_loss: 0.1567 - val\_accuracy: 0.5298  
Epoch 82/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1567 - accuracy: 0.5142 - val\_loss: 0.1645 - val\_accuracy: 0.4862

Epoch 83/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1577 - accuracy:  
0.5148 - val\_loss: 0.1613 - val\_accuracy: 0.4991  
Epoch 84/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1579 - accuracy:  
0.5076 - val\_loss: 0.1617 - val\_accuracy: 0.4898  
Epoch 85/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1573 - accuracy:  
0.5049 - val\_loss: 0.1563 - val\_accuracy: 0.5120  
Epoch 86/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1570 - accuracy:  
0.5142 - val\_loss: 0.1548 - val\_accuracy: 0.5274  
Epoch 87/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1557 - accuracy:  
0.5142 - val\_loss: 0.1551 - val\_accuracy: 0.5169  
Epoch 88/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1553 - accuracy:  
0.5168 - val\_loss: 0.1541 - val\_accuracy: 0.5225  
Epoch 89/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1567 - accuracy:  
0.5111 - val\_loss: 0.1596 - val\_accuracy: 0.5095  
Epoch 90/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1558 - accuracy:  
0.5152 - val\_loss: 0.1545 - val\_accuracy: 0.5243  
Epoch 91/100  
77/77 [=====] - 0s 5ms/step - loss: 0.1570 - accuracy:  
0.5121 - val\_loss: 0.1540 - val\_accuracy: 0.5231  
Epoch 92/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1570 - accuracy:  
0.5152 - val\_loss: 0.1529 - val\_accuracy: 0.5243  
Epoch 93/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1565 - accuracy:  
0.5043 - val\_loss: 0.1610 - val\_accuracy: 0.4985  
Epoch 94/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1560 - accuracy:  
0.5101 - val\_loss: 0.1516 - val\_accuracy: 0.5323  
Epoch 95/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1553 - accuracy:  
0.5144 - val\_loss: 0.1631 - val\_accuracy: 0.4948  
Epoch 96/100  
77/77 [=====] - 1s 8ms/step - loss: 0.1588 - accuracy:  
0.4963 - val\_loss: 0.1639 - val\_accuracy: 0.4874  
Epoch 97/100  
77/77 [=====] - 1s 9ms/step - loss: 0.1584 - accuracy:  
0.5006 - val\_loss: 0.1587 - val\_accuracy: 0.5286  
Epoch 98/100  
77/77 [=====] - 0s 6ms/step - loss: 0.1585 - accuracy:  
0.5027 - val\_loss: 0.1567 - val\_accuracy: 0.5120

```
Epoch 99/100
77/77 [=====] - 0s 6ms/step - loss: 0.1558 - accuracy:
0.4986 - val_loss: 0.1739 - val_accuracy: 0.4695
Epoch 100/100
77/77 [=====] - 0s 5ms/step - loss: 0.1572 - accuracy:
0.5043 - val_loss: 0.1520 - val_accuracy: 0.5311
```

```
[49]: model_performance_graphs(history2, "accuracy")
```



```
[50]: train2_acc = mod2.evaluate(x = X_train, y = Y_train)
print()
print ("Train Accuracy = " + str(train2_acc[1]))
print()
test2_acc = mod2.evaluate(x = X_test, y = Y_test)
print()
print ("Loss = " + str(test2_acc[0]))
print ("Test Accuracy = " + str(test2_acc[1]))
```

```
153/153 [=====] - 0s 2ms/step - loss: 0.1566 -
accuracy: 0.5230
```

```
Train Accuracy = 0.522988498210907
```

```
51/51 [=====] - 0s 3ms/step - loss: 0.1520 - accuracy:
0.5311
```

```
Loss = 0.15203529596328735
```

```
Test Accuracy = 0.5310769081115723
```

Unfortunately, the accuracy of our model reduced when utilizing dropout and batch normalization. When the network is small relative to the dataset, regularization is usually unnecessary. If the

model capacity is already low, lowering it further by adding regularization will hurt performance. Additionally, there are some cases where Batch Normalization can actually hurt the performance of a neural network. For example, if the batch size is too small, Batch Normalization can introduce noise into the training process. Additionally, if the network is very deep, Batch Normalization can make it difficult for the network to learn long-range dependencies.

## 6 Conclusion

In conclusion, I first used a basic stochastic gradient descent algorithm with our wine data. We got an accuracy of nearly 85%. When I introduced dropout and batch normalization, I saw a decrease in our accuracy.

This is the best I could obtain without overfitting. The algorithm might perform better with a better distribution among the different scores. An other explanation might be that wine testing is strongly subjective.

Neural network models are a powerful tool for machine learning. They can be used to solve a wide variety of problems, including classification. However, neural network models can be difficult to train and can be prone to overfitting.

Overfitting is a phenomenon in machine learning where a model performs well on the training data but poorly on the test data. This happens when the model learns the noise in the training data instead of the underlying patterns. While there was a disparity in the accuracies of the two different models, there was little to no difference between the training and testing accuracies of each model.

Overall, neural network models are a powerful tool for machine learning. However, it is important to be aware of their limitations and to use them carefully.

### 6.0.1 References

- <https://www.kaggle.com/code/ryanholbrook/a-single-neuron>
- <https://www.kaggle.com/code/davidzarebski/a-neural-net-for-red-wine-quality-estimation>  
-<https://www.kaggle.com/code/rprkh15/red-wine-quality-xgboost-optuna-neural-networks#Prediction-Using-Neural-Networks>